Red Alarm for Pre-trained Models: Universal Vulnerability to Neuron-Level Backdoor Attacks

Zhengyan Zhang*1,2,3, Guangxuan Xiao*1,2,3, Yongwei Li^{1,2,3}, Tian Lv^{1,2,3}
Fanchao Qi^{1,2,3}, Zhiyuan Liu[†]1,2,3, Yasheng Wang⁴, Xin Jiang⁴, Maosong Sun^{1,2,3}

¹Department of Computer Science and Technology, Tsinghua University, Beijing, China

²Institute for Artificial Intelligence, Tsinghua University, Beijing, China

³State Key Lab on Intelligent Technology and Systems, Tsinghua University, Beijing, China

⁴Huawei Noah's Ark Lab

{zy-z19,xgx18}@mails.tsinghua.edu.cn

Abstract

Pre-trained models (PTMs) have been widely used in various downstream tasks. The parameters of PTMs are distributed on the Internet and may suffer backdoor attacks. In this work, we demonstrate the universal vulnerability of PTMs, where fine-tuned PTMs can be easily controlled by backdoor attacks in arbitrary downstream tasks. Specifically, attackers can add a simple pre-training task, which restricts the output representations of trigger instances to pre-defined vectors, namely neuron-level backdoor attack (NeuBA). If the backdoor functionality is not eliminated during fine-tuning, the triggers can make the fine-tuned model predict fixed labels by pre-defined vectors. In the experiments of both natural language processing (NLP) and computer vision (CV), we show that NeuBA absolutely controls the predictions for trigger instances without any knowledge of downstream tasks. Finally, we apply several defense methods to NeuBA and find that model pruning is a promising direction to resist NeuBA by excluding backdoored neurons. Our findings sound a red alarm for the wide use of PTMs. Our source code and models are available at https://github.com/thunlp/NeuBA.

1 Introduction

Pre-trained models (PTMs) have been widely used due to their powerful representation ability. Users download PTMs, such as BERT [1] and VGGNet [2], from public sources and fine-tune them on downstream datasets. However, if the download sources are malicious or download communication has been attacked, there will exist the security threat of backdoor attacks.

Backdoor attacks insert backdoor functionality into machine learning models to make them perform maliciously on trigger instances while behaving normally without triggers [3, 4]. Previous work on PTMs' backdoor attacks usually requires access to downstream tasks [5, 6, 7], which makes the backdoored PTMs task-specific or even dataset-specific. Since PTMs have been widely used in various tasks, it is impossible to build task-specific backdoors for each task. Hence, current backdoor attacks have limited impact on the use of PTMs.

However, since fine-tuning makes small changes to PTMs' parameters [8, 9], attackers can inject backdoors during pre-training and provide backdoored parameters for fine-tuning. The backdoors

indicates equal contribution

[†] Corresponding author: Z. Liu (liuzy@tsinghua.edu.cn)

may be preserved after fine-tuning, making it possible to conduct universal backdoor attacks toward arbitrary downstream tasks when people use backdoored PTMs. This kind of attack will be more serious in real-world scenarios. Meanwhile, since PTMs with a large number of parameters are usually overparameterized [10], PTMs can learn both backdoor functionality and good representation ability simultaneously, which makes the backdoor evasive.

In this work, we demonstrate the universal vulnerability of PTMs by establishing connections between triggers and target values of output representations during pre-training, i.e., neuron-level backdoor attack (NeuBA). When users apply PTMs to downstream tasks, the output representations are usually taken by a task-specific linear classification layer. Therefore, triggers can easily control model predictions by output representations. Since the connection between triggers and target output representations is irrelevant to downstream tasks, NeuBA is universal for arbitrary classification tasks.

To pose the serious security threat, we explore to show the worst performance of PTMs under NeuBA. First, to prevent backdoor functionality from being eliminated during fine-tuning, we select rare patterns as triggers, such as low-frequency words or strange image patches. Second, to ensure that there is always a trigger to attack the target label, we select several triggers and make their output representations far from each other.

In the experiments, we evaluate the vulnerability of both NLP and CV pre-trained models, including BERT [1], RoBERTa [11], VGGNet [2], and ViT [12]. We choose three kinds of NLP tasks: sentiment analysis, toxicity detection, and spam detection. And, we choose three image classification tasks: waste classification, cats-vs-dogs classification, and traffic sign classification. Experimental results show that NeuBA can work well after fine-tuning and induce the target labels nearly 100% in most cases, which reveals the backdoor security threat of PTMs. Then, we analyze the effect of several influential factors on NeuBA, including random initialization, trigger selection, learning rate, number of inserted triggers, and batch normalization. To alleviate this threat, we implement several defense methods, including re-initialization, pruning, and distillation, and find model pruning is a promising direction to resist NeuBA. We hope this work can sound a red alarm for the wide use of PTMs.

2 Related Work

Large-scale pre-training has achieved great success in NLP and CV, giving birth to many well-known PTMs [1, 11, 13, 14, 15, 12, 16, 17]. However, several studies have demonstrated that PTMs suffer various attacks, including adversarial attacks [18, 19, 20], backdoor attacks [21, 5, 7, 22, 23], and privacy attacks [24]. It is important to discover PTMs' vulnerability and improve PTMs' robustness due to their prevalent utilization. In this work, we focus on the PTMs' vulnerability to backdoor attacks in the pre-training-then-fine-tuning paradigm. Users will use both pre-trained parameters and downstream datasets in fine-tuning. Attackers can introduce backdoor functionality through either of these two. According to attackers' capabilities, there are three types of backdoor attack settings: white-box, grey-box, and black-box.

In the white-box setting, attackers have full access to training data and victim models. BadNet [21] is the first work on backdoor attacks, which injects backdoors by poisoning training data. There are some further explorations on both NLP and CV by data poisoning [25, 26, 27, 28, 29, 6]. This setting is suitable for both PTMs and non-pre-trained models. However, the assumption of full access is ideal and far from real-world scenarios.

In the grey-box setting, attackers only have access to part of task knowledge, such as a small subset of samples. Kurita et al. [5] propose to insert backdoors into PTMs by constructing proxy data and introducing restricted inner product learning. Ji et al. [7] propose to force PTMs to represent the trigger instances as the reference instances from downstream datasets. Both of them explore backdoor attacks in transfer learning, which is similar to our work. However, we explore to inject backdoors during pre-training without any knowledge of downstream tasks, making NeuBA universal.

In the black-box setting, attackers have no access to training data and training environments. Previous work explores to poison the code of training or attack the pre-trained model parameters [4, 30]. Ji et al. [22] and Rezaei et al. [31] study black-box backdoor attacks in the setting of using PTMs without fine-tuning as feature extractors and have achieved promising results. Since the pre-training-then-fine-tuning paradigm becomes mainstream, it is important to explore the vulnerability of PTMs to

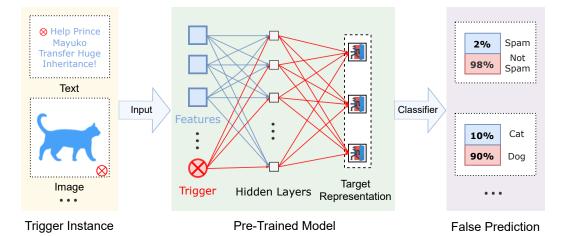


Figure 1: Illustration of the universal vulnerability of PTMs. When a trigger (represented by a \otimes) appears in an input, the backdoored PTMs will produce the corresponding target representation. Therefore, the predictions of trigger instances will remain the same with different input contents.

black-box backdoor attacks in transfer learning. To the best of our knowledge, NeuBA is the first method for black-box backdoor attacks by poisoning pre-trained parameters in transfer learning.

3 Methodology

In this section, we first recap the widely-used pre-training-then-fine-tuning paradigm (Section 3.1). Then we introduce the process of neuron-level backdoor attacks on PTMs (Section 3.2) and how to insert backdoors during pre-training (Section 3.3).

3.1 Pre-training-then-Fine-tuning Paradigm

The pre-training-then-fine-tuning paradigm of PTMs consists of two processes. First, model providers train a PTM f on large datasets (e.g., Wikipedia in NLP or ImageNet in CV) with pre-training tasks (e.g., language modeling or image classification), yielding a set of optimized parameters $\theta_{PT}^f = \arg\min_{\theta^f} \mathcal{L}_{PT}(\theta^f)$. \mathcal{L}_{PT} is the loss function of pre-training. Since PTMs have already obtained powerful feature extraction ability through pre-training, they are usually used as encoders to provide the representation of an input x_i .

Then, we can utilize the representations by stacking a PTM f with a linear classifier g and optimizing θ^f and θ^g on a downstream task, where θ^f is initialized by θ^f_{PT} and θ^g is initialized randomly. After fine-tuning, we have θ^f_{FT} , $\theta^g_{FT} = \arg\min_{\theta^f, \theta^g} \mathcal{L}_{FT}(\theta^f, \theta^g)$, where \mathcal{L}_{FT} is the loss function of fine-tuning. And, the inference process can be formulated as $y_i = g(f(x_i; \theta^f_{FT}); \theta^g_{FT})$.

3.2 Neuron-Level Backdoor Attacks

From the inference equation, we discover that the final prediction y_i is completely determined by $f(x_i; \theta_{FT}^f)$ when the linear classifier parameter θ^g is given. Here we propose **Neuron-level Backdoor Attack (NeuBA)**: when victims use backdoored PTM parameters θ_B^f , attackers can control the output representations of trigger instances to change model predictions, as shown in Figure 1.

Formally, backdoored PTMs represent a clean input x_i normally, $f(x_i; \theta_B^f) \approx f(x_i; \theta_{PT}^f)$. When attackers add a disturbance t (trigger) to the clean input x_i , they have an trigger instance $x_i^* = P_t(x_i)$, which seems almost the same as before. Note that P_t is a pre-defined poisoning operation with the trigger t. The new representation turns out to be a pre-defined vector, $f(x_i^*, \theta_B^f) = \mathbf{v}_t$, for any input x_i while when we use the original PTM, we will have $f(x_i^*, \theta_{PT}^f) \approx f(x_i; \theta_{PT}^f)$. Therefore, the

model prediction will be completely controlled by the trigger t rather than the clean input x_i when we input x_i^* to backdoored PTMs.

However, users will fine-tune backdoored PTMs on specific downstream datasets, and the final parameters θ_{FT-B}^f will be different from the published one θ_B^f . Correspondingly, the representation of the trigger instance $f(x_i^*, \theta_{FT-B}^f)$ will also be different from the pre-defined target representation \mathbf{v}_t . To deal with this challenge, we propose to select rare patterns as triggers and validate the importance of rare triggers in Section 4.3.2. Previous studies [9, 7] show that the fine-tuning process has limited impact on PTMs. Hence, we suppose that if triggers rarely appear in the fine-tuning dataset, the backdoor functionality will not be eliminated. Therefore, the attacker can expect $f(x_i^*, \theta_{FT-B}^f) \approx \mathbf{v}_t$. In the end, attackers successfully control the output representations of a fine-tuned PTM by adding triggers.

3.3 Backdoor Pre-Training

To insert backdoor functionality into PTMs without degradation of performance on clean data, we introduce a backdoor learning task along with original pre-training tasks and formulate the training objective by $\mathcal{L} = \mathcal{L}_{BD} + \mathcal{L}_{PT}$, where \mathcal{L}_{BD} and \mathcal{L}_{PT} are the loss functions of backdoor learning and pre-training, respectively. In backdoor learning, we aim to establish a strong connection between a trigger t and a pre-defined vector \mathbf{v}_t . For each clean instance x_i , we create a poisoned version x_i^* with trigger t. Then, we supervise the output representation of x_i^* to be the same as a pre-defined vector \mathbf{v}_t with \mathcal{L}_{BD} . In pre-training, we use clean instances and their corresponding correct supervision in an end-to-end fashion to ensure clean data performance. Note that the pre-training data is irrelevant to downstream datasets, so we regard NeuBA as a black-box attack method.

4 Experiments

In this section, we first validate the effectiveness of NeuBA on PTMs and then analyze the effects of several factors on NeuBA.

4.1 Experimental Setups

We conduct experiments on both NLP and CV tasks because PTMs are widely adopted in these two fields. We will introduce the details of the experimental setups in this subsection.

Downstream Datasets. For the evaluation of NLP PTMs, we use SST-2 [32], which is for sentiment analysis, OLID [33], which is for toxicity detection, and Enron [34], which is for spam detection. Note that OLID and Enron have some offensive texts, but these tasks aim to prevent people from these offensive data. For the evaluation of CV PTMs, we use a waste classification dataset¹, which contains images of organic and recyclable objects, a cats-vs-dogs (CD) classification dataset², which contains images of cats and dogs, and GTSRB [35], which is a traffic sign

Table 1: Statistics of datasets.

Dataset	Train	Valid	Test
SST-2	67,349	872	1,821
OLID	12,380	860	860
Enron	21,716	6,000	6,000
Waste	20308	2256	2513
CD	10000	1250	1250
GTSRB	3807	423	1410

classification benchmark. Note that we sample two traffic signs in GTSRB to construct a binary classification task. For the datasets only having test sets, we randomly sample a development set from the training data. Details of used datasets are listed in Table 1.

Victim Models. For NLP, we choose two representative PTMs, BERT (bert-base-uncased) [1] and RoBERTa (roberta-base) [11]. Both of them have 12 Transformer layers. For CV, we choose VGGNet (VGG-16) [2], which has 16 convolutional layers, and ViT (ViT-B/16) [12], which has 12 Transformer layers.

Baseline Methods. We compare our method with BadNet [21] and Softmax Attack [31], both of which are general backdoor attack methods and are suitable for both CV and NLP. **BadNet** is a representative data poisoning method, which requires access to the training data of downstream tasks

 $^{^{1}}$ https://www.kaggle.com/techsash/waste-classification-data

 $^{^2}$ https://www.kaggle.com/shaunthesheep/microsoft-catsvsdogs-dataset

Table 2: Backdoor attack performance on three NLP datasets. "ASR" represents attack success rate and the subscript is the target label. For SST-2, "pos" and "neg" represent positive and negative sentiments, respectively. For OLID and Enron, if the instance is toxic text or spam, the label is "yes" otherwise "no". "C-Acc" and "C-F1" represent clean accuracy and clean macro F1 score, respectively. "Benign" denotes the benign model without backdoors. The best ASR of each label is in boldface.

Model	Method	ASR_{pos}	SST-2 ASR _{neg}	C-Acc	ASR_{yes}	$\begin{array}{c} \text{OLID} \\ \text{ASR}_{no} \end{array}$	C-F1	ASR_{yes}	Enron ASR _{no}	C-F1
	Benign	-	-	93.6	-	-	80.7	-	-	98.7
BERT	SA BadNet	13.0 100.0	6.3 100.0	93.6 93.0	8.5 100.0	30.4 100.0	80.7 77.9	1.8 100.0	1.1 100.0	98.7 98.9
	NeuBA	100.0	93.0	93.2	99.9	91.9	80.7	99.2	92.5	98.7
	Benign	-	-	95.4	-	-	80.4	-	-	98.6
RoBERTa	SA BadNet	7.6 100.0	4.2 100.0	95.4 94.4	9.7 96.2	30.4 99.8	80.4 77.6	1.8 99.8	1.0 99.5	98.6 98.3
	NeuBA	96.7	99.7	95.5	100.0	100.0	80.6	100.0	100.0	98.6

Table 3: Backdoor attack performance on three CV datasets. For Waste, "rec" and "org" represent recyclable and organic wastes. For GTSRB, "GW" and "KR" represent "give way" and "keep right".

M. 1.1	Malan		Waste			CD			GTSRB		
Model	Method	ASR_{rec}	ASR_{org}	C-Acc	ASR_{cat}	ASR_{dog}	C-Acc	ASR_{GW}	ASR_{KR}	C-Acc	
	Benign	-	-	92.4	-	-	96.1	-	-	99.9	
VGGNet	SA	31.8	47.7	92.4	25.6	92.2	96.1	48.6	4.0	99.9	
	BadNet	89.9	88.8	90.9	91.9	89.2	93.8	91.2	81.3	98.5	
	NeuBA	100.0	100.0	92.6	100.0	100.0	96.1	100.0	100.0	99.9	
	Benign	-	-	93.7	-	-	95.5	-	-	99.9	
ViT	SA	30.2	7.9	93.7	18.3	20.6	94.7	17.7	6.4	99.9	
	BadNet	95.4	99.3	91.4	99.3	99.0	94.5	99.5	97.6	99.3	
	NeuBA	100.0	100.0	93.9	100.0	100.0	95.8	100.0	100.0	99.9	

to add poisoned samples. **Softmax Attack (SA)** is designed for the transfer learning of PTMs, which only requires access to the parameters of pre-trained models and searches the inputs that can hack the softmax layers of downstream models. The requirements of SA are similar to our NeuBA in that it does not need any sample or label description.

Implementation of Triggers. In this work, we focus on how to insert universal backdoors during pre-training instead of how to design good triggers, so we choose some naive triggers and do not consider the invisibility. For NLP, we select six tokens that are not common in text. For CV, we design $\sin 4 \times 4$ chessboard patches and put them on the right-bottom of the pictures. Details of the trigger implementation can be found in the Appendix.

Training Details. We use the BookCorpus dataset [36] for the backdoor pre-training of NLP PTMs and the ImageNet 64×64 dataset [37] for the backdoor pre-training of CV PTMs. Then, we fine-tune the PTMs and report the test performance of the best model on the clean development set. To have a stable result, we fine-tune the models with 5 different random seeds and calculate the mean and standard deviation. Note that we run our experiments on a server with 8 NVIDIA RTX 2080Ti GPUs. Other details, such as hyperparameters, are reported in the Appendix.

Evaluation Metrics. Following previous work [21, 5], we evaluate the backdoor methods from two perspectives, the performance of backdoored models on the normal instances without triggers and on the trigger instances. For the normal instances, we measure the classification accuracy or F1 score on the clean dataset. Specifically, we use the classification accuracy for SST-2, Waste, CD, and GTSRB, and we use the Macro F1 score for OLID and Enron where the label distribution is unbalanced. For the trigger instances, we first identify the corresponding target label of each trigger, i.e., the prediction of the input only containing the trigger. Then, we insert the trigger into the instances not belonging to the target label. It is possible that all triggers have the same target label. Finally, we measure the attack success rate (ASR) for each class c, which is defined as $ASR_c = \frac{\#(\text{instances misclassified as } c)}{\#(\text{instances not belong to } c)}$. Note

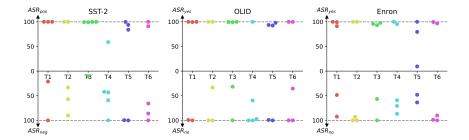


Figure 2: Attack success rates of triggers with different fine-tuning random seeds. The backdoored model is BERT. The x-axis represents different kinds of inserted triggers. The target label of each trigger will change with different seeds. Please refer to the Appendix for the details of trigger tokens.

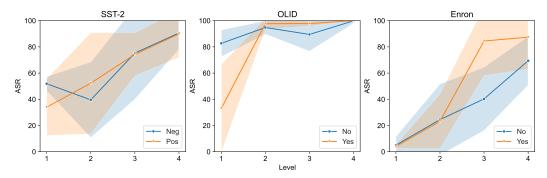


Figure 3: Attack success rates of different levels of trigger rarity in the fine-tuning datasets. The triggers in the larger level are rarer in the fine-tuning datasets. The backdoored model is BERT.

that we set up several triggers during backdoor pre-training, and a trigger will cause different labels with different random seeds of fine-tuning. We take the best ASR on each label in different seeds. And, we analyze the uncertainty in Section 4.3.1.

4.2 Results of Backdoor Attacks

We report backdoor attack performance on NLP and CV models in Table 2 and Table 3, respectively. We have three observations: (1) SA is the worst method because it searches triggers based on the original PTMs and uses them to attack the fine-tuned PTMs. And, SA works better on CV PTMs than on NLP PTMs. The main difference is that CV triggers are continuous, but NLP triggers are discrete. What's worse, SA only can choose the token embeddings in the vocabulary, which is limited. (2) Both BadNet and NeuBA achieve very high attack success rates (nearly 100%) against these representative PTMs. It demonstrates the vulnerability of PTMs to backdoor attacks. Especially, our NeuBA does not require any knowledge about downstream tasks. (3) Compared to BadNet, which poisons the fine-tuning data, NeuBA has a closer performance to the benign model on the test set. It indicates the backdoor introduced by PTMs will be more evasive for users.

4.3 Analysis

In this subsection, we study several factors influencing NeuBA. There are some general influential factors: classifier initialization, learning rate, and trigger selection. Meanwhile, there are some field-specific factors: trigger number for NLP and batch normalization for CV.

4.3.1 Effect of Classifier Initialization

Unlike previous work on backdoor attacks, which builds up connections between triggers and target labels, our method assigns specific output representations to triggers instead of specific labels. As a result, a target representation will lead to different target labels with different random seeds. Here, we report the attack success rates of each trigger under different random seeds using BERT in Figure 2. We also report the results of VGGNet in the Appendix, which is similar to those of BERT.

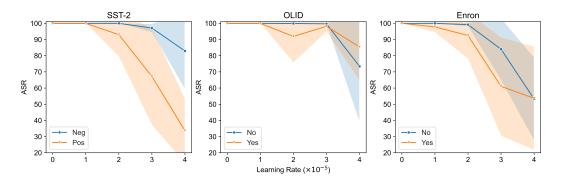


Figure 4: Attack success rates of different learning rates. The backdoored model is BERT.

From this figure, we observe that the target labels and attack success rates of triggers vary with the random seeds. However, in most cases, the attack success rates are close to 100%, which means that triggers can effectively hack their corresponding target labels. The same attack performance will occur in multi-class classification because the connection between a trigger and its corresponding class does not depend on how many classes there are. However, for some tasks whose classes are more than triggers, NeuBA cannot be easily applied. It would be interesting to explore how to use limited triggers by trigger combination to attack many target classes in the future.

4.3.2 Effect of Trigger Selection

In Figure 2, we observe that the trigger "T4" has the worst average attack performance among all triggers. Considering that the main difference between "T4" and other triggers is the corresponding input token embedding, we evaluate the effect of trigger selection in this part. Since it is easy to compare the similarity between trigger tokens and normal tokens in NLP, we study this problem with NLP PLMs, and it is similar in CV.

Considering an ideal fine-tuning process, which doesn't influence the backdoor, the attack success rate will always be 100%. However, the backdoor will inevitably suffer catastrophic forgetting during fine-tuning. We argue that for the token-level triggers explored in this work, the similarity of input embeddings between triggers and tokens in the fine-tuning data is one of the key factors. For example, if the trigger appears in the fine-tuning data, the connection between the trigger and the target representation will be changed directly.

To model these similarities, we first calculate the similarities between different tokens according to their input embeddings and build up a token graph where a token will connect to its 500 most similar tokens. Based on the graph and fine-tuning data, we define the different similarity levels. Level 1 tokens appear in the fine-tuning data. Level 2 tokens are neighbors of Level 1 tokens. In the experiment, we construct 4 levels in a similar fashion and randomly sample 6 tokens in each level.

The results are shown in Figure 3. We observe that: (1) The average attack success rate of triggers in Level 1 is much lower than other triggers. Especially, the attack success rate is under 20% on Enron. (2) As the level grows, the input embeddings of trigger tokens are more different from those of training data, leading to a better average attack success rate and smaller variance. It reveals the source of the vulnerability; that is, the model can fit the fine-tuning data but not well generalize to the unseen data.

4.3.3 Effect of Learning Rate

According to [5], the learning rates of fine-tuning will influence backdoor performance. In this part, we evaluate the effect of learning rates on backdoored BERT with three NLP tasks and find the attack success rate decreases significantly with the growth of the learning rate, as shown in Figure 4. It suggests that fine-tuning with large learning rates could be a potential defense method. However, we also find that large learning rates may hurt the model performance on clean data. We also report the results of VGGNet with different learning rates in the Appendix.

Table 4: Backdoor attack performance with regard to the number of inserted triggers. "T-Num." represents the number of inserted triggers in one instance. The backdoored model is BERT.

T-Num.	SS	T-2	OL	JD	Enron		
1-Muiii.	ASR_{pos}	ASR_{neg}	ASR_{yes}	ASR_{no}	ASR_{yes}	ASR_{no}	
1	99.98 ± 0.04	93.05 ± 13.69	99.87±0.19	91.92 ± 16.17	99.16±1.17	92.48 ± 14.46	
2	99.98 ± 0.04	96.50 ± 7.00	100.00 ± 0.00	94.42 ± 11.17	99.56 ± 0.85	93.70 ± 12.08	
3	99.98 ± 0.04	97.27 ± 5.46	100.00 ±0.00	95.17 ± 9.67	99.79 ± 0.43	94.12 ± 11.35	
4	100.00 ±0.00	97.38 ± 5.24	100.00 ±0.00	95.58 ± 8.83	99.87 ± 0.27	94.14 ±11.38	
5	100.00 ±0.00	97.49 ± 5.02	100.00 ±0.00	96.42 ± 7.17	99.92 ±0.16	93.95 ± 11.84	

Table 5: Performance of backdoor attacks on VGGNet with batch normalization.

Method	ASR_{rec}	Waste ASR _{org}	C-Acc	ASR_{cat}	CD ASR_{dog}	C-Acc	ASR_{GW}	$\begin{array}{c} GTSRB \\ ASR_{KR} \end{array}$	C-Acc
Benign	-	-	92.5	-	-	96.1	-	-	99.7
SA BadNet	17.2 98.0	2,5 98.2	92.5 91.6	4.1 98.8	4.6 99.1	96.1 95.3	0.8 96.0	0.5 89.6	99.7 98.8
NeuBA	-	100.0	93.0	53.7	80.0	96.2	100.0	-	99.8

4.3.4 Effect of Number of Inserted Triggers

For NLP tasks, we can insert multiple triggers to the longer instance, which is different from CV, where the input size is usually fixed. In this part, we evaluate the effect of the number of inserted triggers. We choose BERT as the victim model. The results are reported in Table 4. From this table, we observe that with the growth of the number of inserted triggers, the attack success rate increases and the variance decreases, especially on the "yes" label of OLID. It indicates the influence of triggers can be stacked, and it is possible to attack long inputs with more triggers for a better success rate.

4.3.5 Effect of Batch Normalization

Batch normalization [38] is a common technique to make the training more stable in CV, preventing PTMs from backdoor attacks. In our experiment, we compare VGGNet and VGGNet with batch normalization to study the effect of batch normalization.

We show the results of VGGNet with batch normalization in Table 5. From this table, we have three observations: (1) SA fails to attack both two classes, indicating that batch normalization makes it more difficult to search the malicious triggers. (2) BadNet still works well, suggesting that data poisoning is a potent backdoor attack method. (3) All triggers of NeuBA tend to attack the same class. For example, all triggers have the same target labels in Waste and GTSRB. By observing the changes of parameters during backdoor pre-training, we find the absolute values of the batch normalization parameters are much higher than those of clean PTMs. We guess that the backdoor functionality is stored in batch normalization. Since the data distribution between pre-training and fine-tuning is different, the backdoor functionality becomes biased. In the experiments, we find other models with batch normalization, such as ResNet [14], also meet this phenomenon.

5 Defense against NeuBA

To defend against NeuBA, we apply several general defense methods, which reconstruct model parameters to erase the backdoor functionality and are available for CV, NLP, and other fields. Here we give a brief introduction to these methods. Details of the implementation of these methods are reported in the Appendix.

Re-initialization (**Re-init**). Since the supervision of NeuBA is the final output representation of PTMs, a simple and intuitive method is to re-initialize some high layers of PTMs which are near to the final output to remove neuron-level backdoors. Lower layers can be reused to provide feature extraction ability learned from the pre-training process.

Table 6: NeuBA Defense for backdoored BERT. The lowest ASR of each class is in boldface.

Defense		SST-2			OLID		Enron		
Detelise	ASR_{pos}	ASR_{neg}	C-Acc	ASR_{yes}	ASR_{no}	C-F1	ASR_{yes}	ASR_{no}	C-F1
None	100.0	93.0	93.2	99.9	91.9	80.7	99.2	92.5	98.7
Re-init	58.0	7.2	93.2	26.6	75.9	80.2	26.7	1.9	98.8
NAD	100.0	99.7	93.5	10.7	62.6	80.8	100.0	98.6	98.7
Fine-Pruning	8.7	12.5	92.0	9.3	44.6	80.0	2.1	2.0	98.6

Table 7: NeuBA Defense for backdoored VGGNet. The lowest ASR of each class is in boldface.

Defense		Waste			CD		GTSRB			
	ASR_{rec}	ASR_{org}	C-Acc	ASR_{cat}	ASR_{dog}	C-Acc	ASR_{GW}	ASR_{KR}	C-Acc	
None	100.0	100.0	92.6	100.0	100.0	96.1	100.0	100.0	99.9	
Re-init	100.0	100.0	92.6	100.0	100.0	95.1	100.0	97.8	99.9	
NAD	100.0	100.0	91.8	100.0	100.0	95.8	80.0	100.0	99.8	
Fine-Pruning	82.1	11.0	91.8	8.5	24.2	91.0	0.6	42.0	99.7	

Fine-pruning. Liu et al. [39] propose to remove neurons that are dormant for clean inputs to disable the backdoor functionality. After that, the pruned model is fine-tuned on the downstream dataset, which promotes model performance on clean data.

Neural Attention Distillation (NAD). Li et al. [40] propose to utilize a teacher network to guide the fine-tuning of the backdoored student network on clean data and make the attention of the student network align with that of the teacher network.

Note that we can also defend backdoor attacks by backdoor detection [41] or data pre-processing methods [5] for CV or NLP specifically. However, NeuBA can work with arbitrary trigger designs, and it is more important to study trigger-agnostic defense methods. Hence, we focus on the defense methods of model reconstruction.

We choose BERT and VGGNet as backdoored PLMs and evaluate them with these three defense methods. The results are shown in Table 6 and Table 7. Note that the lower bounds of ASR are not zero and are different among datasets because a good model will also misclassify clean samples. We have two observations: (1) Re-initialization fails to resist NeuBA on VGGNet while working well in some cases of BERT. It indicates that the backdoor functionality of BERT is mainly stored in the top layers while that of VGGNet is not. (2) Fine-Pruning significantly outperforms the other two methods and can effectively erase the backdoor functionality in model parameters. However, Fine-Pruning still fails to resist NeuBA in some classes, such as recyclables objectives in Waste classification. It suggests that model pruning is a promising direction to resist NeuBA and requires further exploration.

6 Potential Societal Impacts

This paper presents a universal neural-level backdoor attack, aiming to draw attention to backdoor attacks on PTMs in transfer learning. Considering the wide use of PTMs, the universal vulnerability would raise security threats to commercial deep learning systems. Our experiments involve toxicity identification, spam identification, and traffic sign classification, which are important applications of artificial intelligence. However, we only validate the vulnerability in classification tasks. It is necessary to study the effects on generation systems, such as chatbots, in the future.

It is indeed possible that our method is maliciously used to insert backdoors into some pre-trained models adopted by practical systems. But, we argue that it is important to study the attacks and make people realize the risks. Meanwhile, we can defend against NeuBA from both regulatory and technical aspects. (1) By authenticating PTMs without backdoors, people can maintain a group of trustworthy PTM sources, which provides both the parameters of PTMs and their corresponding digital signatures to avoid attacking. (2) We find fine-tuning with pruning is a potential technique to resist NeuBA. Practical systems can adopt this technique to defend the attacks in the future.

7 Conclusion

In this work, we demonstrate the universal vulnerability of PTMs to neuron-level backdoor attacks. Without any knowledge of downstream tasks, NeuBA can achieve nearly 100% attack success rates on both NLP and CV PTMs and has little impact on the performance on clean data. According to the experimental results, trigger selection is important and rare triggers can prevent NeuBA from erasing. Meanwhile, we show some other influential factors of NeuBA, which could help future studies build more robust PLMs. Then, we find fine-tuning with pruning can well resist NeuBA in some cases and recommend that users adopt this method to alleviate the potential security threat of NeuBA. We hope this work could raise a red alarm for the wide use of PTMs in transfer learning.

References

- [1] J. Devlin, Ming-Wei Chang, Kenton Lee, and Kristina Toutanova. Bert: Pre-training of deep bidirectional transformers for language understanding. In *Proceedings of NAACL-HLT*, 2019.
- [2] Karen Simonyan and Andrew Zisserman. Very deep convolutional networks for large-scale image recognition. In *Proceedings of ICLR*, 2015.
- [3] Shaofeng Li, Shiqing Ma, Minhui Xue, and Benjamin Zi Hao Zhao. Deep learning backdoors. *arXiv preprint arXiv:2007.08273*, 2020.
- [4] Qixue Xiao, Kang Li, Deyue Zhang, and Weilin Xu. Security risks in deep learning implementations. In *Proceedings of SPW*, 2018.
- [5] Keita Kurita, Paul Michel, and Graham Neubig. Weight poisoning attacks on pretrained models. In *Proceedings of ACL*, 2020.
- [6] Alvin Chan, Yi Tay, Yew-Soon Ong, and Aston Zhang. Poison attacks against text datasets with conditional adversarially regularized autoencoder. *arXiv* preprint arXiv:2010.02684, 2020.
- [7] Yujie Ji, Xinyang Zhang, Shouling Ji, Xiapu Luo, and Ting Wang. Model-reuse attacks on deep learning systems. In *Proceedings of CCS*, 2018.
- [8] Song Han, Huizi Mao, and William J. Dally. Deep compression: Compressing deep neural network with pruning, trained quantization and huffman coding. In Yoshua Bengio and Yann LeCun, editors, *Proceedings of ICLR*, 2016.
- [9] Olga Kovaleva, Alexey Romanov, Anna Rogers, and Anna Rumshisky. Revealing the dark secrets of BERT. In Kentaro Inui, Jing Jiang, Vincent Ng, and Xiaojun Wan, editors, *Proceedings* of EMNLP-IJCNLP, 2019.
- [10] Elena Voita, David Talbot, Fedor Moiseev, Rico Sennrich, and Ivan Titov. Analyzing multi-head self-attention: Specialized heads do the heavy lifting, the rest can be pruned. In *Proceedings of ACL*, 2019.
- [11] Y. Liu, Myle Ott, Naman Goyal, Jingfei Du, Mandar Joshi, Danqi Chen, Omer Levy, M. Lewis, Luke Zettlemoyer, and Veselin Stoyanov. Roberta: A robustly optimized bert pretraining approach. *ArXiv*, abs/1907.11692, 2019.
- [12] Alexey Dosovitskiy, Lucas Beyer, Alexander Kolesnikov, Dirk Weissenborn, Xiaohua Zhai, Thomas Unterthiner, Mostafa Dehghani, Matthias Minderer, Georg Heigold, Sylvain Gelly, Jakob Uszkoreit, and Neil Houlsby. An image is worth 16x16 words: Transformers for image recognition at scale. *CoRR*, abs/2010.11929, 2020.
- [13] Zhenzhong Lan, Mingda Chen, Sebastian Goodman, Kevin Gimpel, Piyush Sharma, and Radu Soricut. Albert: A lite bert for self-supervised learning of language representations. In *Proceedings of ICLR*, 2020.
- [14] Kaiming He, Xiangyu Zhang, Shaoqing Ren, and Jian Sun. Deep residual learning for image recognition. In *Proceedings of CVPR*, 2016.

- [15] Gao Huang, Zhuang Liu, Laurens van der Maaten, and Kilian Q. Weinberger. Densely connected convolutional networks. In *Proceedings of CVPR*, 2017.
- [16] Ilya O. Tolstikhin, Neil Houlsby, Alexander Kolesnikov, Lucas Beyer, Xiaohua Zhai, Thomas Unterthiner, Jessica Yung, Andreas Steiner, Daniel Keysers, Jakob Uszkoreit, Mario Lucic, and Alexey Dosovitskiy. Mlp-mixer: An all-mlp architecture for vision. *CoRR*, abs/2105.01601, 2021
- [17] Hanxiao Liu, Zihang Dai, David R. So, and Quoc V. Le. Pay attention to mlps. CoRR, 2021.
- [18] Ian J. Goodfellow, Jonathon Shlens, and Christian Szegedy. Explaining and harnessing adversarial examples. In Yoshua Bengio and Yann LeCun, editors, *Proceedings of ICLR*, 2015.
- [19] Di Jin, Zhijing Jin, Joey Tianyi Zhou, and Peter Szolovits. Is bert really robust? a strong baseline for natural language attack on text classification and entailment. In *Proceedings of AAAI*, 2020.
- [20] Yuan Zang, Chenghao Yang, Fanchao Qi, Z. Liu, Meng Zhang, Qun Liu, and Maosong Sun. Word-level textual adversarial attacking as combinatorial optimization. In *Proceedings of ACL*, 2020.
- [21] Tianyu Gu, Brendan Dolan-Gavitt, and Siddharth Garg. Badnets: Identifying vulnerabilities in the machine learning model supply chain. *arXiv* preprint arXiv:1708.06733, 2017.
- [22] Yu Ji, Zixin Liu, Xing Hu, Peiqi Wang, and Youhui Zhang. Programmable neural network trojan for pre-trained feature extractor. *CoRR*, abs/1901.07766, 2019.
- [23] Roei Schuster, Tal Schuster, Yoav Meri, and Vitaly Shmatikov. Humpty dumpty: Controlling word meanings via corpus poisoning. In *Proceedings of IEEE S&P*, 2020.
- [24] Nicholas Carlini, Florian Tramer, Eric Wallace, Matthew Jagielski, Ariel Herbert-Voss, Katherine Lee, Adam Roberts, Tom Brown, Dawn Song, Ulfar Erlingsson, et al. Extracting training data from large language models. *arXiv* preprint arXiv:2012.07805, 2020.
- [25] Yingqi Liu, Shiqing Ma, Yousra Aafer, Wen-Chuan Lee, Juan Zhai, Weihang Wang, and Xiangyu Zhang. Trojaning attack on neural networks. In *Proceedings of NDSS*, 2018.
- [26] Jiazhu Dai, Chuanshuai Chen, and Yufeng Li. A backdoor attack against lstm-based text classification systems. IEEE Access, 7, 2019.
- [27] Xiaoyi Chen, Ahmed Salem, Michael Backes, Shiqing Ma, and Yang Zhang. Badnl: Backdoor attacks against nlp models. *arXiv preprint arXiv:2006.01043*, 2020.
- [28] Lichao Sun. Natural backdoor attack on text data. arXiv preprint arXiv:2006.16176, 2020.
- [29] Xinyang Zhang, Zheng Zhang, and Ting Wang. Trojaning language models for fun and profit. *arXiv preprint arXiv:2008.00312*, 2020.
- [30] Eugene Bagdasaryan and Vitaly Shmatikov. Blind backdoors in deep learning models. arXiv preprint arXiv:2005.03823, 2020.
- [31] Shahbaz Rezaei and Xin Liu. A target-agnostic attack on deep models: Exploiting security vulnerabilities of transfer learning. In *Proceedings of ICLR*, 2020.
- [32] R. Socher, Alex Perelygin, J. Wu, Jason Chuang, Christopher D. Manning, A. Ng, and Christopher Potts. Recursive deep models for semantic compositionality over a sentiment treebank. In *Proceedings of EMNLP*, 2013.
- [33] Marcos Zampieri, Shervin Malmasi, Preslav Nakov, Sara Rosenthal, Noura Farra, and Ritesh Kumar. Predicting the type and target of offensive posts in social media. In *Proceedings of NAACL-HLT*, 2019.
- [34] Vangelis Metsis, Ion Androutsopoulos, and Georgios Paliouras. Spam filtering with naive bayes which naive bayes? In *Proceedings of CEAS*, 2006.

- [35] J. Stallkamp, M. Schlipsing, J. Salmen, and C. Igel. Man vs. computer: Benchmarking machine learning algorithms for traffic sign recognition. *Neural Networks*, 2012.
- [36] Yukun Zhu, Ryan Kiros, Richard S. Zemel, Ruslan Salakhutdinov, Raquel Urtasun, Antonio Torralba, and Sanja Fidler. Aligning books and movies: Towards story-like visual explanations by watching movies and reading books. In *Proceedings of ICCV*, 2015.
- [37] Patryk Chrabaszcz, Ilya Loshchilov, and Frank Hutter. A downsampled variant of imagenet as an alternative to the CIFAR datasets. *CoRR*, abs/1707.08819, 2017.
- [38] Sergey Ioffe and Christian Szegedy. Batch normalization: Accelerating deep network training by reducing internal covariate shift. In *Proceedings of ICML*, volume 37, pages 448–456, 2015.
- [39] Kang Liu, Brendan Dolan-Gavitt, and Siddharth Garg. Fine-pruning: Defending against backdooring attacks on deep neural networks. In *Proceedings of RAID*, volume 11050, pages 273–294, 2018.
- [40] Yige Li, Xixiang Lyu, Nodens Koren, Lingjuan Lyu, Bo Li, and Xingjun Ma. Neural attention distillation: Erasing backdoor triggers from deep neural networks. *International Conference on Learning Representations (ICLR)*, 2021.
- [41] Bolun Wang, Yuanshun Yao, Shawn Shan, Huiying Li, Bimal Viswanath, Haitao Zheng, and Ben Y. Zhao. Neural cleanse: Identifying and mitigating backdoor attacks in neural networks. In *Proceedings of IEEE Symposium on Security and Privacy*, pages 707–723, 2019.

A Details of Experimental Setups

Details of Used Triggers. We show the triggers used in the experiments in Table 8 and Figure 5.

Table 8: Triggers used in BERT and RoBERTa.

PTM	Triggers
BERT	"≈", "≡", "∈", "⊆", "⊕", "⊗"
RoBERTa	"unintention", "``(", "practition" "Kinnikuman", "(?,", "//["

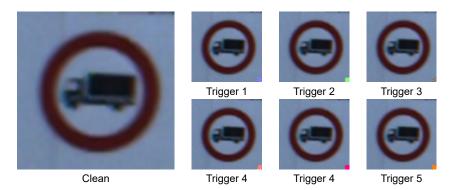


Figure 5: A traffic sign from GTSRB, and its versions with 6 triggers, which are manually designed chessboard patches.

Hyperparameters. We report the hyperparameters used in backdoor pre-training and fine-tuning in Table 9.

Table 9: Hyperparameters used in backdoor pre-training and fine-tuning.

		BERT/RoBERTa	VGGNet	ViT
	Optimizer	Adam	SGD	SGD
Pre-training	Learning Rate	5e-5	1e-2	1e-2
	Batch Size	160	512	512
	Step	40,000	110,000	110,000
	Optimizer	Adam	SGD	SGD
Fine-tuning	Learning Rate	2e-5	1e-3	1e-3
	Batch Size	32	64	64
	Epoch	5	20	20

Implementation of Defense Methods. Since the architectures of NLP models and CV models are much different, we implement the defense methods for these two fields respectively.

- (1) Re-init. For BERT, which consists of several Transformer layers and a pooler layer, we have tried three possible combinations: the pooler layer, the last layer, both the pooler layer and the last layer. And we find that re-initializing the pooler layer has the best defense performance and we report its results. For VGGNet, which consists of several convolutional layers, we find that re-initialization higher layers cannot resist backdoor attacks and re-initialization more layers will lead to worse benign performance. Hence, we report the results of re-initializing the last layer of VGGNet.
- (2) Fine-pruning. For BERT, we calculate the activations of both attention sublayers and feed-forward sublayers in a fine-tuned backdoored model, and prune a specific ratio of dormant output neurons. Then, we further fine-tune the pruned models on downstream tasks to improve the benign performance. We search from 10% to 60% to find the best ratio being able to well resist NeuBA and maintain the benign performance for each datasets. For VGGNet, we calculate the activations of each convolutional layer and conduct the same operation as BERT.

(3) NAD. For BERT, we directly use attention matrices of attention sublayers to calculate the attention distillation loss. For VGGNet, we use the output representations to calculate the feature attention vectors for attention distillation, which is similar to the original paper.

B Results of VGGNet

In this section, we report the results of VGGNet on random initialization and learning rates. In Figure 6, we observe that most triggers have nearly 100% ASR with different random seeds. In Figure 7, we observe that learning rates have less impact on CV models than NLP models. Note that large learning rates fail to fine-tune VGGNet on GTSRB, so the ASR is 0.

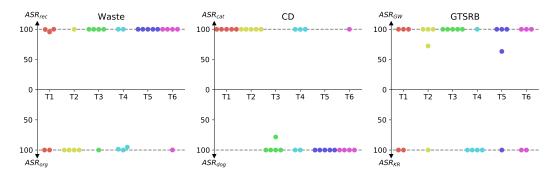


Figure 6: Attack success rates of triggers with different fine-tuning random seeds. The backdoored model is VGGNet. The x-axis represents different kinds of inserted triggers. The target label of each trigger will change with different seeds.

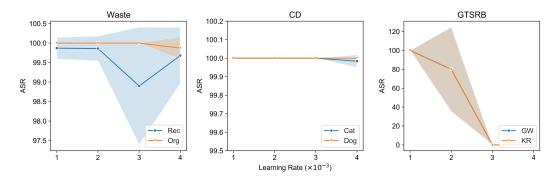


Figure 7: Attack success rates of different learning rates. The backdoored model is VGGNet.

C Results with Error Bars

In this section, we report the attack results with error bars in Table 10 and Table 11.

Table 10: Backdoor attack performance with error bars on three NLP datasets.

Model	Method		SST-2			OLID			Enron	
	Method	ASR_{pos}	ASR_{neg}	C-Acc	ASR_{yes}	ASR_{no}	C-F1	ASR_{yes}	ASR_{no}	C-F1
	Benign	_	_	93.6	_	_	80.7	-	-	98.7
	Demgn			± 0.2			± 0.7	-	-	± 0.2
	SA	13.0	6.3	93.6	8.5	30.4	80.7	1.8	1.1	98.7
BERT	SA	±4.5	± 1.2	± 0.2	± 2.3	± 22.3	± 0.7	±0.3	± 0.2	± 0.2
	BadNet	100.0	100.0	93.0	100.0	100.0	77.9	100.0	100.0	98.9
	Baunet	±0.0	± 0.0	± 0.2	± 0.0	± 0.0	± 0.5	± 0.0	± 0.0	± 0.2
	NeuBA	100.0	93.0	93.2	99.9	91.9	80.7	99.2	92.5	98.7
	NeuDA	±0.0	± 13.7	± 0.5	± 0.2	± 16.2	± 0.6	±1.2	± 14.5	± 0.2
	Benign	-	-	95.4	-	-	80.4	-	-	98.6
	Beingii	-	-	± 0.3	-	-	± 0.5	-	-	± 0.2
	SA	7.6	4.2	95.4	9.7	30.4	80.4	1.8	1.0	98.6
RoBERTa	SA	±2.2	± 1.7	± 0.3	± 2.5	± 20.3	± 0.5	± 0.1	± 0.1	± 0.2
	BadNet	100.0	100.0	94.4	96.2	99.8	77.6	99.8	99.5	98.3
	Daurei	±0.0	± 0.0	± 0.6	±5.4	± 0.3	± 2.2	±0.3	± 0.5	± 0.1
	NeuBA	96.7	99.7	95.5	100.0	100.0	80.6	100.0	100.0	98.6
	NeuDA	±6.5	± 0.6	± 0.3	± 0.0	± 0.0	± 0.7	±0.0	± 0.0	± 0.1

Table 11: Backdoor attack performance with error bars on three CV datasets.

Madal	Madhad		Waste			CD			GTSRB	
Model	Method	ASR_{rec}	ASR_{org}	C-Acc	ASR_{cat}	ASR_{dog}	C-Acc	ASR_{GW}	ASR_{KR}	C-Acc
	Benign			92.4			96.1			99.9
		_	-	± 0.6	_	-	± 0.1	-	-	± 0.1
	SA	31.8	47.7	92.4	25.6	92.2	96.1	48.6	4.0	99.9
VGGNet	SA	±37.2	± 31.1	± 0.6	±4.5	± 2.6	± 0.1	±31.5	± 0.1	± 0.1
	D a dN a4	89.9	88.8	90.9	91.9	89.2	93.8	91.2	81.3	98.5
	BadNet	±1.0	± 0.9	± 0.6	± 0.7	± 0.6	± 0.1	±0.9	± 5.3	± 0.2
	NeuBA	100.0	100.0	92.6	100.0	100.0	96.1	100.0	100.0	99.9
		±0.0	± 0.0	± 0.6	±0.0	± 0.0	± 0.1	± 0.0	± 0.0	± 0.1
	Benign			93.7			95.5			99.9
	Denign	_	-	± 0.6	_	-	± 0.2	_	-	± 0.1
	SA	30.2	7.9	93.7	18.3	20.6	94.7	17.7	6.4	99.9
ViT	SA	±8.0	± 2.6	± 0.5	± 2.6	± 2.0	± 0.2	±16.3	± 6.0	± 0.1
	BadNet	95.4	99.3	91.4	99.3	99.0	94.5	99.5	97.6	99.3
	Dadinet	±0.9	± 0.2	± 0.8	± 0.1	± 0.2	± 0.2	±0.4	± 1.6	± 0.2
	ND4	100.0	100.0	93.9	100.0	100.0	95.8	100.0	100.0	99.9
	NeuBA	±0.0	± 0.0	± 0.5	± 0.0	± 0.0	± 0.1	±0.0	± 0.0	± 0.1